Do Synthetic Personas Predict Real Audience Response? A Headline-Ranking Study with a Generic-Reader Baseline Historical Results and a Partial Corrected Reanalysis

Alexandre Cristovão Maiorano
alexandre@lumytics.com

Abstract

Do demographic personas improve an LLM's prediction of audience engagement? We compare a fixed ten-persona panel with a generic-reader baseline on Upworthy headlines, with additional news and Reddit comparisons. This revision corrects metric identity, selection, ties, and claims in the original release. Historical Upworthy estimates favor the baseline: Goodman–Kruskal $\gamma = 0.361$ versus 0.084, and top-1 49.2% versus 34.6%. These used a post-selection one-sided winner screen and first-index tie breaking; they remain historical results, not corrected Kendall estimates. The baseline does not dominate every dataset: Reddit top-1 and captured lift favor the panel. A retrospective GPT-4.1 diagnostic uses a two-sided, multiplicity-adjusted winner screen, explicit metric denominators, and expected credit for tied choices. On 189 paired packages, its baseline advantage persists. This diagnostic reuses recorded post-parser values, retains unequal sampling budgets, and cannot validate original responses. Historical Gemini observations were unavailable for equivalent reanalysis. The evidence supports a baseline comparison for the tested configurations, not a general causal claim about persona harm. Public reports support numerical verification; complete historical reproduction from raw observations remains incomplete.

1 Introduction

LLM-based audience simulation offers an inexpensive way to compare messages before live testing. Demographic conditioning can approximate aspects of survey samples [1], but a copy-ranking system must predict measured engagement. We study this narrower question with Upworthy headline experiments [16].

Revision and evidence status. An external artifact review prompted a full consistency audit of release `7c1ab46`. The tables reproduce stored reports, but that agreement does not validate the estimators. The original rank implementation computed Goodman–Kruskal gamma while the paper called it Kendall tau. The winner screen selected the observed best arm before a one-sided test. Undefined coefficients were excluded without effective denominators. Claims about baseline superiority everywhere, prompt robustness, and a causal explanation also exceeded the artifacts.

This revision distinguishes historical tables, retaining original numbers with corrected labels, from an offline GPT-4.1 reanalysis. The latter reuses recovered checkpoint scores after replay of the pinned original producer reproduces the published report. It changes selection, metrics, and tie credit, but cannot undo historical response parsing. No new provider generations were collected for this revision.

We ask (RQ1) how this panel and generic baseline compare within tested packages, (RQ2) how agreement varies with the observed outcome signal, and (RQ3) what the dataset, aggregation, seed, and model comparisons establish. This evaluates a specified harness rather than all persona methods or audience representativeness.

2 Related Work

Profile-conditioned models can approximate some aggregate survey distributions [1]; broader computational social science applications raise validity questions [24]. Interview-grounded agents reproduce General Social Survey responses at approximately 85% of participants' own test–retest consistency in the original version of Park et al. [18]. That result uses rich individual grounding and a different outcome from clicks.

Other work documents option-order sensitivity [5], limited effects of persona attributes on subjective annotation [10], persona-related degradation on objective tasks [23], caricature [2], and biased outcomes from generated personas [15]. Reasoning biases in persona-assigned models are also documented [8].

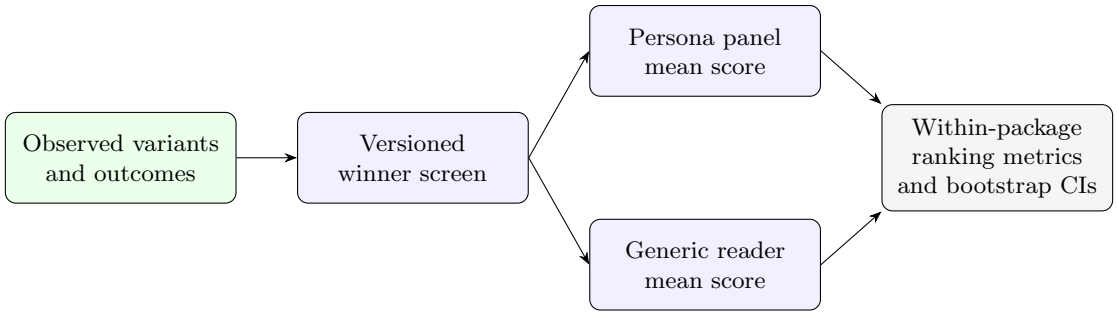


Figure 1: Comparison protocol. Outcomes determine selection; the two scoring arms are compared within selected packages. Historical outputs use gamma and the legacy screen. Corrected derivatives additionally report tau-b and adjusted selection. Selection does not establish that excluded packages are unpredictable.

These motivate a baseline, but do not explain this study's gap.

Upworthy [16] supports engagement analyses including negativity effects [19]. LOLA [22] studies LLM-assisted content experiments; differences in selection, models, and evaluation prevent a direct numerical comparison here. MIND [21] and Reddit reposts [14] are additional observational settings, not interchangeable headline experiments. Construct validity [4] distinguishes intent from behavior. Gamma [7] must also be distinguished from Kendall correlation and its treatment of ties [11, 12].

3 Sim-to-Real Validity Framework

A package contains variants with observed outcome rates and predicted scores (Figure 1). Upworthy uses randomized traffic, although grouping distinct headlines can combine image variants. MIND compares different articles within a subcategory. Reddit groups reposts of the same image in the same subreddit; timing and exposure remain uncontrolled. Within-package ranking does not establish absolute CTR calibration or a causal copy effect in observational data.

Metric identity. For concordant and discordant pair counts C, D, the historical implementation used

$$\gamma = (C - D)/(C + D),$$

discarding pairs tied in either variable [7]. It is undefined when $C + D = 0$ and is not Kendall τ_b. The corrected code separately computes

$$\tau_b = \frac{C - D}{\sqrt{(C + D + T_x)(C + D + T_y)}},$$

where T_x, T_y count pairs tied only in the respective variable [12]. Spearman uses average ranks [20]. Constant rank inputs yield undefined values.

Corrected top-1 and bottom-1 use uniform expected credit: $|S_{\text{pred}} \cap S_{\text{real}}|/|S_{\text{pred}}|$ for the corresponding extreme sets. Captured lift averages observed rates across tied predicted maxima, then scales distance from the observed minimum by the observed range. Constant outcomes make lift undefined. Historical top-1 selected the first maximum and remains ordering-dependent.

4 Method Under Test

The study harness scores each headline with either a demographic persona or a generic reader. It operates independently of any commercial product. The fixed ten-persona JSON takes inspiration from Upworthy-era U.S. social-media demographics [9]. These descriptors do not specify a deterministic behavioral-profile generator or establish representativeness.

Elicitation and aggregation. The prompt requests JSON with a `click_intent` number in $[0, 1]$ (Appendix A). The primary collector averages recorded observations:

$$\widehat{s}(v) = \frac{1}{|O_v|} \sum_{(i,d) \in O_v} y_{i,d}(v).$$

With complete draws this equals the mean over personas and draws. With missing draws it can differ from the historical H5 mean, which weights draw-collapsed personas equally. H5 already evaluated alternative aggregations; they were not wholly left to future work.

Sampling and parsing. The main Gemini configuration uses temperature 0.8, three draws per persona

and thirty generic baseline draws per variant. GPT-4.1 instead uses three draws per persona versus ten baseline draws. Gemini forwards generation seeds; the GPT runner sends none. Seeds alone cannot guarantee identical provider responses over time.

The legacy parser could extract a first number, clamp values, and accept booleans. Invalid parsing could terminate without retry. Historical event frequencies are unknown. The revised collector accepts only a finite, non-boolean JSON number in range under the exact requested key. Invalid parses are recorded as terminal missing observations. Configuration hashes and a new protocol namespace prevent mixing new checkpoints with legacy scores.

5 Experimental Setup

Corpus. Upworthy rows are grouped by experiment and distinct headline, aggregating identical text. Packages require at least two headlines and 3,000 impressions per headline; sampling uncertainty matters in online experiments [13]. The locked exploratory corpus has 6978 variants in 1695 packages. Original manifests remain unchanged.

Historical winner screen. The original procedure chose the observed best variant and compared it with the observed runner-up using a one-sided pooled two-proportion z test at $p < 0.05$, selecting 399 packages. The direction was chosen from the data: even a two-arm null comparison can pass in either direction, so the nominal one-sided threshold does not imply a five-percent selection error. Comparing only the runner-up need not separate a winner from every opponent. Outcome selection changes the evaluation population. Packages failing a screen are not thereby unpredictable.

Corrected screen. The revised rule requires a unique observed maximum and compares it with every opponent using two-sided pooled z tests. For k variants the comparison family contains $m = k(k-1)/2$ unordered pairs. All winner–opponent comparisons must pass α/m, with $\alpha = 0.05$. The screening statistic is $\min(1, m \max_{j \neq w} p_{wj})$. Bonferroni's inequality motivates the adjustment [17]. Error control is approximate and per package, conditional on independent binomial observations and adequate normal approximation. It is not an across-package false-discovery guarantee or proof of a population winner. Repeated exposures and Reddit voting further limit these assumptions. In the restored corpus 190 packages pass, all within the historical selected set. A stricter screen does not demonstrate that removed packages had false winners.

Estimation and denominators. Package metrics are averaged without size weighting. Corrected reports include $n_{\text{total}}, n_{\text{valid}}, n_{\text{undefined}}$ and metric-specific package IDs. Each paired gap uses packages valid in both arms for that metric. Empty estimates are undefined, not zero. Percentile bootstrap CIs use 10,000 package resamples, seed 42 [6]. They condition on selection and recorded scores and omit generation, selection, and outcome-measurement uncertainty. Historical coefficient-specific valid counts are not fully recorded; coverage totals are not those denominators.

Planning and robustness scope. The archived power calculation is an earlier planning artifact. Fisher-z power for one correlation does not justify an average of within-package coefficients, and Cohen's conventions [3] do not automatically calibrate gamma. This revision makes no prospective power or equivalence claim. Time-stamped preregistration of all thresholds is not established.

Historical robustness artifacts cover three Upworthy splits, MIND [21], Reddit [14], three Gemini models, seed runs, and aggregation. No identifiable result artifact supports the previous prompt-robustness claim. Reddit's manifest records many more selected packages than the evaluated subset; its subset-selection procedure is undocumented.

6 Results

Historical Upworthy comparison (RQ1). Table 1 and Figure 2 retain original estimates. The generic baseline has higher gamma (0.361 versus 0.084) and top-1 (49.2% versus 34.6%). These are conditional descriptions under the legacy screen and tie rule. Marginal CIs are not a paired test, and gamma is not assigned a calibrated medium-effect interpretation.

Figure 3 shows the historical screening statistic, not calibrated winner probabilities. The archived all-package panel gamma is 0.054 (CI [0.025, 0.083]), also retaining legacy metric and parsing semantics.

Signal strength (RQ2). The historical association between the legacy signal statistic and top-1 is Spearman 0.059 (CI [0.011, 0.107]) over 1678 packages. It has not been recomputed for Gemini under corrected screening and tie handling.

Table 1: Historical Upworthy exploratory estimates under the legacy winner screen. Panel $n_{total} = 399$; baseline $n_{total} = 398$; correlation-specific valid counts are unavailable. The stored rank coefficient is Goodman–Kruskal gamma, not Kendall tau. Top-1 uses first-index tie breaking. Random top-1 reference 30.4%. Marginal 95% bootstrap CIs; this table is not a corrected reanalysis.

Method	Top-1 [95% CI]	Goodman–Kruskal γ [95% CI]	Spearman ρ [95% CI]
No-persona baseline	49.2% [44.2, 54.3]	0.361 [0.289, 0.430]	0.350 [0.281, 0.417]
Persona panel	34.6% [29.8, 39.1]	0.084 [0.020, 0.147]	0.104 [0.035, 0.172]

Table 2: Historical estimates under the legacy screen and first-index tie breaking. n_{total} is panel/baseline coverage, not the correlation-specific valid count (unavailable). Lift is the fraction of the observed outcome-rate range captured by the selected variant (panel/baseline). The panel exceeds the baseline on Reddit lift and top-1 (discussed in text); near-zero gamma does not establish equivalence or unpredictability. Upworthy rows use one draw per persona; Table 1 uses multiple draws.

Dataset	n_{total} p/b	Persona γ	Baseline γ	Gap	Lift p/b
Upworthy exploratory (2013–15)	399/398	0.079	0.361	+0.282	0.45/0.58
Upworthy holdout (2013–15)	402/401	0.100	0.398	+0.298	0.48/0.56
Upworthy confirmatory (2013–15)	1866/1864	0.070	0.393	+0.323	0.44/0.55
MIND news (2019)	74/74	0.177	0.214	+0.038	0.47/0.51
Reddit titles (2008–12)	600/599	-0.009	-0.026	−0.017	0.45/0.33

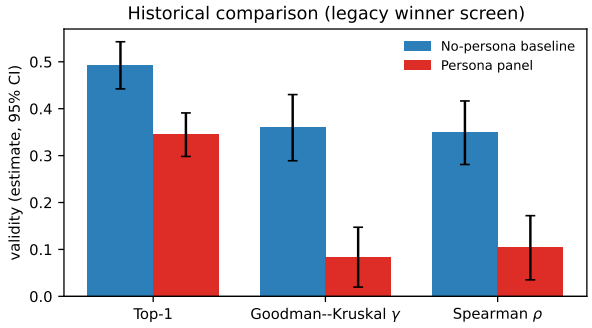


Figure 2: Historical comparison under the legacy screen. Coefficient-specific valid counts are unavailable. The rank coefficient is gamma, not a corrected Kendall estimate.

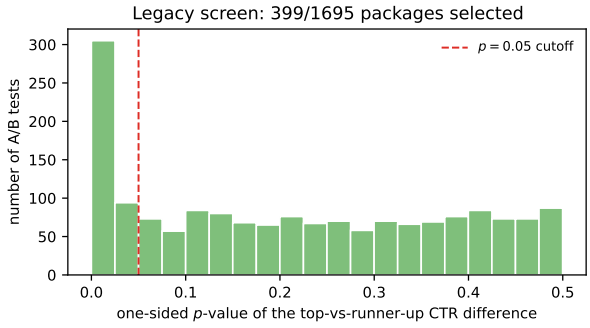


Figure 3: Historical one-sided top-versus-runner-up screen selecting 399 packages. Selection of direction from observed rates invalidates a literal one-sided error interpretation. Nonsignificant packages need not be unpredictable.

Dataset comparisons (RQ3). Table 2 shows higher baseline gamma on the three Upworthy splits and a smaller difference on MIND, where packages compare articles. On Reddit, the panel has higher top-1 (40.3% versus 30.2%) and lift (0.452 versus 0.333). The original baseline-superiority claim for every dataset and metric was false. Reddit gamma CIs include zero; this establishes neither equivalence nor absence of a title-related signal. Posting time and visibility were not tested as explanations. Subset-selection and coverage limitations further restrict interpretation.

Seed, model, ties, and aggregation. The archived three-seed SDs (0.025 for gamma; 0.008 for top-1) use the population convention over those seeds, not sample SDs or precise general-variability estimates. The baseline has higher gamma on three tested Gemini tiers (Table 6, Appendix A); this cannot establish universal model or prompt robustness.

The tie report rounds scores before counting constant packages; the rank code uses exact equality. That count cannot recover excluded gamma values. Undefined coefficients cannot be silently assigned zero, and the reports do not justify a conservative-lower-

Table 3: Historical H5 aggregation comparison (n_{total} = 399 for panel rules). The baseline has 398 covered packages; correlation-specific valid counts are unavailable. The mean row weights draw-collapsed personas equally, whereas the primary collector weights recorded observations equally. These estimands can differ when draws are missing; the stored reports do not identify the cause of their top-1 discrepancy. No tested rule reaches the baseline's reported gamma. This descriptive comparison does not identify a failure mechanism. Legacy screen, first-index tie breaking, 95% marginal bootstrap CIs.

Aggregation	Goodman–Kruskal γ [95% CI]	Top-1 [95% CI]
No-persona baseline	0.361 [0.289, 0.430]	49.2% [44.2, 54.3]
Equal-persona mean	0.084 [0.020, 0.147]	34.8% [30.1, 39.3]
Panel median	0.141 [0.076, 0.205]	39.3% [34.6, 44.1]
Trimmed mean (20%)	0.081 [0.017, 0.144]	34.3% [29.6, 38.8]
Per-persona z-mean	0.064 [0.001, 0.128]	33.6% [29.1, 38.3]
Borda (rank fusion)	0.081 [0.018, 0.146]	35.3% [30.8, 40.1]
Condorcet (pairwise)	0.079 [0.014, 0.146]	35.8% [31.3, 40.6]

bound claim for the gap.

Table 3 preserves six historical aggregation rules, all with lower reported gamma than the baseline. Equal-persona and observation-weighted means can differ with missing draws. The reports do not establish why their top-1 results differ, or whether bias, sampling noise, scale, or another factor explains relative performance.

GPT-4.1 historical replay and corrected diagnostic. The recovered checkpoint contains scalars, not original responses. Offline replay of the pinned producer reproduces all published cross-family report fields within numerical tolerance. Historical coverage is 399 panel and 395 baseline/paired packages. Thirty panel observations versus ten baseline draws and no generation seed make it an unequal-budget comparison.

Table 4 applies the corrected screen and estimators to those scores. Coverage becomes 190 panel, 189 baseline, and 189 paired packages. Baseline Kendall τ_b is 0.325 versus 0.049. Paired CIs for the displayed gaps exclude zero, supporting a positive baseline gap in this retrospective diagnostic. It is not independent replication, repaired raw parsing, or a budget-matched estimate. Gamma and tau-b remain separate. Corrected historical Gemini estimates are unavailable pending recovery of its observations.

7 Discussion and Limitations

The comparisons make a generic-reader baseline necessary for this task. A persona panel adds assumptions about audience composition and behavior that require predictive validation. Historical Upworthy reports and the conditional GPT diagnostic favor the baseline; Reddit decision metrics favor the panel. Findings depend on configuration and outcome.

Mechanism remains open. The baseline may use general headline knowledge, while role-play may add unhelpful conditioning. Neither explanation was isolated experimentally. Similar score spreads, selected error examples, and rank-only aggregation failures cannot distinguish bias from noise or identify a mechanism. A balanced, interleaved experiment retaining raw responses would better test those hypotheses. Intent versus engagement remains a construct-validity limitation [4].

Ground truth and generalization. The corrected screen addresses selection under approximate binomial tests; it does not make outcomes error-free. Exclusion changes the target population and relative task difficulty. Failure to reject equality does not imply no predictor can recover an ordering. MIND and Reddit are observational; correlated exposures, time, image, audience mix, and topic may matter. Training contamination and generalization to current audiences or other formats are not excluded by a historical corpus.

Reproduction limits. Aggregate reports allow re-export, not complete verification of their generating process. Gemini observations and some historical report producers were not recovered. The GPT checkpoint supports replay and partial reanalysis, but range-checking stored scalars cannot retroactively validate the permissive parser. New strict-parser checkpoints cannot be mixed with legacy scores. Prompt robustness, deterministic demographic-to-panel construction, full Reddit subset lineage, and preregistration chronology remain evidence gaps.

Preserving historical artifacts and separating corrected derivatives prevents numerical verification from being mistaken for empirical revalidation.

8 Conclusion

This fixed demographic persona panel should be compared with a generic-reader LLM baseline. Historical Upworthy results favor that baseline, whereas Reddit top-1 and lift favor the panel. An offline GPT-4.1

Table 4: Retrospective GPT-4.1 diagnostic on recorded post-parser values. Corrected winner screen and uniform expected tie credit; marginal percentile bootstrap CIs. Gap is baseline minus panel on paired packages; top-1 gap is in percentage points. Counts are per-metric valid panel/baseline/paired. Original responses are unavailable, so this does not validate the historical parser. Sampling remains 30 panel observations versus 10 baseline draws per variant.

Metric	Panel [95% CI]	Baseline [95% CI]	Paired gap [95% CI]	n_{valid} p/b/pair
Goodman–Kruskal γ	0.049 [-0.061, 0.159]	0.325 [0.223, 0.429]	0.270 [0.132, 0.407]	190/189/189
Kendall τ_b	0.049 [-0.061, 0.159]	0.325 [0.222, 0.429]	0.270 [0.132, 0.407]	190/189/189
Expected top-1	39.5% [32.6, 46.3]	60.6% [53.7, 67.7]	20.9 pp [11.1, 30.4]	190/189/189

diagnostic retains a positive paired baseline gap, conditional on recorded post-parser scores and unequal budgets.

The original coefficient was gamma, not Kendall tau. Selection and undefined denominators also required correction. We release distinct metrics, expected tie credit, explicit valid and paired counts, and an adjusted winner screen. These improve the protocol without establishing universal persona harm or a causal mechanism. Complete corrected historical Gemini analysis requires the missing observations.

Data and Code Availability

The package at `https://github.com/alemaiorano/sim-to-real-validity` includes sources, exporters, aggregate historical reports, the fixed panel, and original corpus manifests. Reexporting these reports reproduces their displayed numbers, not their generation from raw observations. Original reports are hash-pinned and unchanged. Separate corrected GPT derivatives, input hashes, and offline replay verification are in `reanalysis/2026-10-31/`. The recovered checkpoint is not redistributed. Independent recomputation requires the exact checkpoint and corpus, and historical Gemini observations remain missing. New generations are new experiments, not substitutes for those observations. `REPLICATION.md` distinguishes these tiers and limitations.

Ethics and Conflict of Interest

The study uses public archive data and synthetic personas; no new human subjects were recruited. The author has a commercial interest in copy-simulation tools. This evaluates a study harness rather than a specific product. Disclosure, archived adverse results, explicit limitations, and released artifacts support scrutiny; an unverified preregistration chronology is not presented as a safeguard.

AI Tools Disclosure

AI coding assistants were used for code, analysis, literature checks, and manuscript revision. The historical manuscript disclosed Claude and GPT-family assistance. This correction also used Codex and an Astra high subagent to challenge the correction plan. These tools do not supply independent empirical validation; the author remains responsible for the work. The evaluated models are the historical Gemini configurations and GPT-4.1 recorded in the reports. No paid model experiments were run for the correction.

References

[1] Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, 2023. doi: 10.1017/pan.2023.2.

[2] Myra Cheng, Tiziano Piccardi, and Diyi Yang. CoMPosT: Characterizing and evaluating caricature in LLM simulations. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, 2023. doi: 10.18653/v1/2023.emnlp-main.669.

[3] Jacob Cohen. *Statistical Power Analysis for the Behavioral Sciences*. Routledge, 2nd edition, 1988. doi: 10.4324/9780203771587.

[4] Lee J. Cronbach and Paul E. Meehl. Construct validity in psychological tests. *Psychological Bulletin*, 52(4):281–302, 1955. doi: 10.1037/h0040957.

[5] Ricardo Dominguez-Olmedo, Moritz Hardt, and Celestine Mendler-Dünner. Questioning the survey responses of large language models. In *Advances in Neural Information Processing Systems*,

pages 45850–45878, 2024. doi: 10.52202/079017-1458.

[6] B. Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979. doi: 10.1214/aos/1176344552.

[7] Leo A. Goodman and William H. Kruskal. Measures of association for cross classifications. *Journal of the American Statistical Association*, 49 (268):732–764, 1954. doi: 10.1080/01621459.1954.10501231.

[8] Shashank Gupta, Vaishnavi Shrivastava, Ameet Deshpande, Ashwin Kalyan, Peter Clark, Ashish Sabharwal, and Tushar Khot. Bias runs deep: Implicit reasoning biases in persona-assigned LLMs. In *International Conference on Learning Representations (ICLR)*, 2024.

[9] Jesse Holcomb, Jeffrey Gottfried, and Amy Mitchell. News use across social media platforms. Technical report, Pew Research Center, 11 2013. URL `https://www.pewresearch.org/journalism/2013/11/14/news-use-across-social-media-platforms/`. In collaboration with the John S. and James L. Knight Foundation. Accessed 2026-06-03.

[10] Tiancheng Hu and Nigel Collier. Quantifying the persona effect in LLM simulations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10289–10307, 2024. doi: 10.18653/v1/2024.acl-long.554.

[11] M. G. Kendall. A new measure of rank correlation. *Biometrika*, 30(1/2):81–93, 1938. doi: 10.1093/biomet/30.1-2.81.

[12] M. G. Kendall. The treatment of ties in ranking problems. *Biometrika*, 33(3):239–251, 1945. doi: 10.1093/biomet/33.3.239.

[13] Ron Kohavi, Roger Longbotham, Dan Sommerfield, and Randal M. Henne. Controlled experiments on the web: survey and practical guide. *Data Mining and Knowledge Discovery*, 18(1): 140–181, 2009. doi: 10.1007/s10618-008-0114-1.

[14] Himabindu Lakkaraju, Julian McAuley, and Jure Leskovec. What's in a name? understanding the interplay between titles, content, and communities in social media. In *Proceedings of the Seventh International AAAI Conference on Weblogs and Social Media*, pages 311–320, 2013. doi: 10.1609/icwsm.v7i1.14408. 2013 proceedings; publisher portal migration date is 2021.

[15] Ang Li, Haozhe Chen, Hongseok Namkoong, and Tianyi Peng. LLM Generated Persona is a Promise with a Catch, 2025.

[16] J. Nathan Matias, Kevin Munger, Marianne Aubin Le Quere, and Charles Ebersole. The upworthy research archive, a time series of 32,487 experiments in U.S. media. *Scientific Data*, 8(1): 195, 2021. doi: 10.1038/s41597-021-00934-7.

[17] NIST/SEMATECH. e-handbook of statistical methods: Bonferroni's method. `https://www.itl.nist.gov/div898/handbook/prc/section4/prc473.htm`. Accessed 2026-10-31; used for the multiplicity principle, not the binomial approximation.

[18] Joon Sung Park, Carolyn Q. Zou, Aaron Shaw, Benjamin Mako Hill, Carrie Cai, Meredith Ringel Morris, Robb Willer, Percy Liang, and Michael S. Bernstein. Generative agent simulations of 1,000 people, 2024. URL `https://arxiv.org/abs/2411.10109v1`.

[19] Claire E. Robertson, Nicolas Pröllochs, Kaoru Schwarzenegger, Philip Pärnamets, Jay J. Van Bavel, and Stefan Feuerriegel. Negativity drives online news consumption. *Nature Human Behaviour*, 7(5):812–822, 2023. doi: 10.1038/s41562-023-01538-4.

[20] C. Spearman. The proof and measurement of association between two things. *The American Journal of Psychology*, 15(1):72–101, 1904. doi: 10.2307/1412159.

[21] Fangzhao Wu, Ying Qiao, Jiun-Hung Chen, Chuhan Wu, Tao Qi, Jianxun Lian, Danyang Liu, Xing Xie, Jianfeng Gao, Winnie Wu, and Ming Zhou. MIND: A large-scale dataset for news recommendation. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 3597–3606, 2020. doi: 10.18653/v1/2020.acl-main.331.

[22] Zikun Ye, Hema Yoganarasimhan, and Yufeng Zheng. LOLA: LLM-assisted online learning algorithm for content experiments. *arXiv preprint arXiv:2406.02611*, 2024.

[23] Mingqian Zheng, Jiaxin Pei, Lajanugen Logeswaran, Moontae Lee, and David Jurgens. When "a helpful assistant" is not really helpful: Personas in system prompts do not improve performances of large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 15126–15154, 2024. doi: 10.18653/v1/2024.findings-emnlp.888.

[24] Caleb Ziems, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. Can large language models transform computational social science? *Computational Linguistics*, 50(1):237–291, 2024. doi: 10.1162/coli_a_00502.

A Persona Panel, Prompt, and Reproducibility Details

Fixed panel. The ten-persona JSON is the reproducible input. Demographic inspiration [9] is not a deterministic recipe for recreating behavioral descriptions. The unweighted panel is not an estimated population sample.

Historical first-person template.

> You are simulating a specific person reacting to a news headline on social media in 2014.
> Person: {persona description} (age {band}, {gender}, mainly on {platform}, news engagement: {level}).
> Headline: {headline}.
> As this exact person, how likely are you to click this headline? Respond ONLY with JSON: {"click_intent": ¡number between 0 and 1¿}.

The exact templates, including quotation marks and the first-person, third-person, and baseline variants, are authoritative in the versioned collector. A prompt implementation alone does not establish a completed experiment.

Coverage and checkpoints. Each arm requires a score for every variant, not successful completion of every planned draw. The primary mean averages remaining observations. New collectors record terminal invalid/missing outputs and bind checkpoints to protocol and configuration hashes. Transport errors have bounded retries; invalid parsing is not retried. Gemini forwards draw-specific seeds; GPT has no seed parameter.

Historical pilot and model comparison. Tables 5 and 6 retain the pilot and cross-tier gap. Overlap of pilot CIs does not establish equivalence; the cheapest model was selected. This numerical pilot supersedes untraceable statements about a different preliminary pilot.

Illustrative errors and panel. Table 7 shows selected cases where the baseline found the observed winner and the panel did not; they are not representative causal evidence. Table 8 presents the fixed panel descriptions. Full fields and grounding notes remain in `data/persona_panel.json`.

Table 5: Historical pilot on a common subset ($n = 100$ packages). The intervals overlap; this is not an equivalence test. The cheapest model was selected. 95% bootstrap CIs.

Model	Top-1	Goodman–Kruskal γ	Spearman ρ
`gemini-2.5-flash`	0.390 [0.29, 0.49]	0.116 [-0.01, 0.24]	0.133 [-0.01, 0.27]
`gemini-3.1-flash-lite`	0.330 [0.24, 0.42]	0.035 [-0.10, 0.17]	0.051 [-0.09, 0.19]
`gemini-3.5-flash`	0.390 [0.29, 0.49]	0.122 [-0.01, 0.25]	0.139 [-0.00, 0.28]

Table 6: Historical cross-model comparison on a common subset ($n_{total} = 100$ per condition). The baseline has higher reported gamma on these three tested Gemini models. This does not establish a model-independent effect; correlation-specific valid counts are unavailable.

Model	Persona γ	Baseline γ	Gap
`gemini-3.1-flash-lite`	0.035	0.401	+0.367
`gemini-2.5-flash`	0.116	0.219	+0.103
`gemini-3.5-flash`	0.122	0.404	+0.282

Table 7: Examples of real winning headlines that the no-persona baseline ranked first but the persona panel did not (of 126 such cases). These are selected illustrations, not evidence of a systematic causal mechanism.

Real winning headline	CTR	#var
Bookmark This For When You're Sad And Instantly Be 100 Times Happier	1.38%	4
10 Amazing International Photos That Will Give You Some Perspective	1.92%	4
If You Saw Someone Bullying A Kid For Being Gay Would You Speak Out...	1.11%	2
If You Watch Hilary Swank Annoy A Whiny Senator, $1 Will Go To Pun...	1.67%	3
Seven Wonderful People Share Secrets That Are Killing Them. This C...	1.99%	4
A 9-Year-Old Asks McDonald's CEO A Simple Question. She's Not Lovin...	3.23%	4
Did These Words Really Come Out Of Conan's Mouth?	4.45%	3
We Use This Word So Often We Probably Don't Realize We're Saying It...	1.68%	4
Did You Know You Can Go For A Swim In Google Earth?	0.88%	3
Cookie Monster Is In Jail? I'm Actually Really Happy About The Reas...	2.30%	4

Table 8: The full fixed persona panel (P=10), inspired by Upworthy-era audience demographics [9], with the verbatim behavioural descriptions used in the prompt (for replicability). News engagement and political lean are in the replication package (`data/persona_panel.json`).

ID	Demographics	Behavioural description (verbatim)
p01	18-29, F, Facebook	29-year-old woman, scrolls Facebook on her phone during breaks; clicks human-interest and feel-good stories shared by friends; rarely seeks news directly.
p02	18-29, M, Twitter	24-year-old man, heavy Twitter user, follows journalists and culture accounts; clicks surprising or contrarian takes; low tolerance for clickbait he's seen before.
p03	30-49, F, Facebook	38-year-old working mother, checks Facebook in the evening; drawn to headlines about family, health, and local impact; shares uplifting or outrage-inducing stories.
p04	30-49, M, Reddit	41-year-old man, browses Reddit and Facebook; skeptical of sensational headlines, clicks when a headline promises concrete information or a strong narrative payoff.
p05	50-64, F, Facebook	57-year-old woman, mostly on Facebook; clicks emotionally resonant stories and anything about social causes she follows; less drawn to internet-slang headlines.
p06	50-64, M, Facebook	60-year-old man, reads news on Facebook and cable; clicks political and economic headlines, distrusts overtly hyped 'you won't believe' framing.
p07	18-29, F, YouTube	22-year-old college student, video-first, on YouTube and Instagram; clicks visually evocative or identity-relevant headlines; very sensitive to authenticity.
p08	30-49, M, Facebook	45-year-old man, casual Facebook user; clicks curiosity-gap headlines that promise a quick interesting fact, especially maps, lists, and 'what happened next' stories.
p09	65+, F, Facebook	68-year-old retiree, active on Facebook with family; clicks heartwarming, nostalgic, or cause-driven stories; avoids headlines that feel like tricks.
p10	30-49, F, Twitter	34-year-old professional, follows news and advocacy on Twitter; clicks data-driven or justice-oriented angles; fatigued by repetitive viral formats.